



Intro to seq2seq

Машинное обучение в диалоговых системах
2 апреля 2019

Cheskova E - cheskidova@phystech.edu

Outline

1. Seq2seq architectures with rnn
 - a. Applications in NLP
 - b. Simple encoder, simple decoder
 - c. Decoder with attention
2. Transformer
 - a. Encoder
 - i. Multi-head self-attention
 - ii. LayerNorm & residual connections
 - iii. Position-wise feed-forward
 - iv. Positional Encoding
 - b. Decoder
 - i. Multi-head attention with encoder outputs
 - ii. Masking
3. Optional Part

Outline

1. Seq2seq architectures with rnn

a. Applications in NLP

- b. Simple encoder, simple decoder
- c. Decoder with attention

2. Transformer

a. Encoder

- i. Multi-head self-attention
- ii. LayerNorm & residual connections
- iii. Position-wise feed-forward
- iv. Positional Encoding

b. Decoder

- i. Multi-head attention with encoder outputs
- ii. Masking

Machine translation

PARALLEL CORPUS



Generative chatbot

INPUT: *Hi!*

OUTPUT: *Hi!*

INPUT: *How are you?*

OUTPUT: *I am good. Thanks for asking.*

INPUT: *You're asking me out. That's so cute. What's your name again?*

OUTPUT: *Forget it.*

INPUT: *No, no, it's my fault—we didn't have a proper introduction*

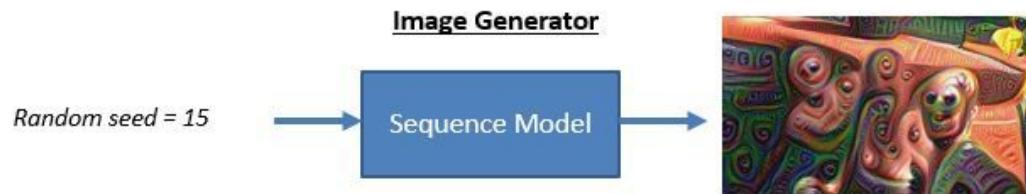
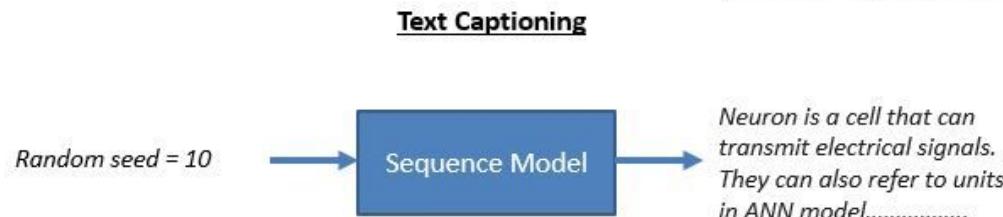
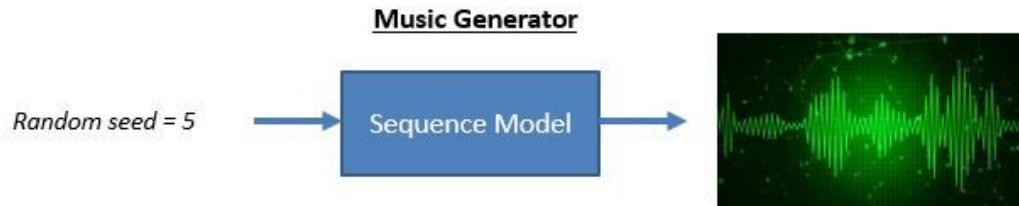
OUTPUT: *Cameron.*

INPUT: *The thing is, Cameron—I'm at the mercy of a particularly hideous breed of loser. My sister. I can't date until she does.*

OUTPUT: *Seems like she could get a date easy enough...*



Other applications



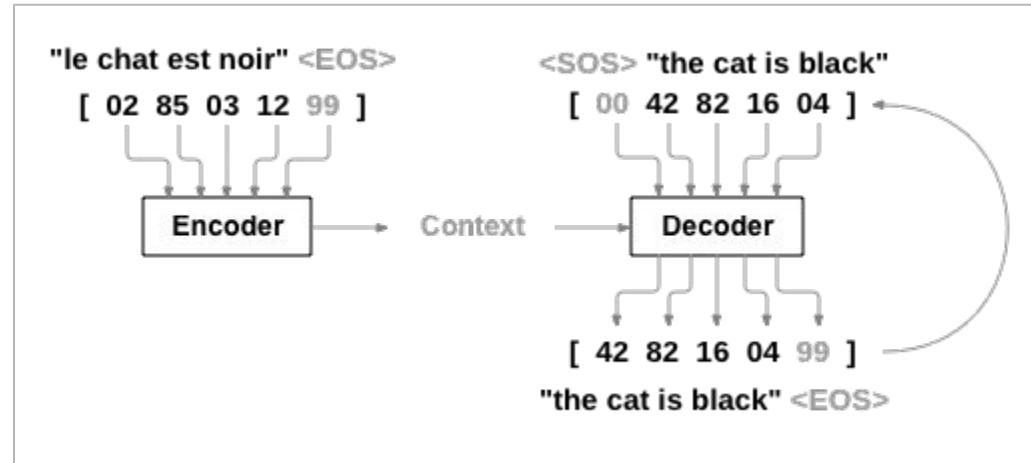
Outline

1. Seq2seq architectures with rnn
 - a. Applications
 - b. Simple encoder, simple decoder**
 - c. Decoder with attention
2. Transformer
 - a. Encoder
 - i. Multi-head self-attention
 - ii. LayerNorm & residual connections
 - iii. Position-wise feed-forward
 - iv. Positional Encoding
 - b. Decoder
 - i. Multi-head attention with encoder outputs
 - ii. Masking

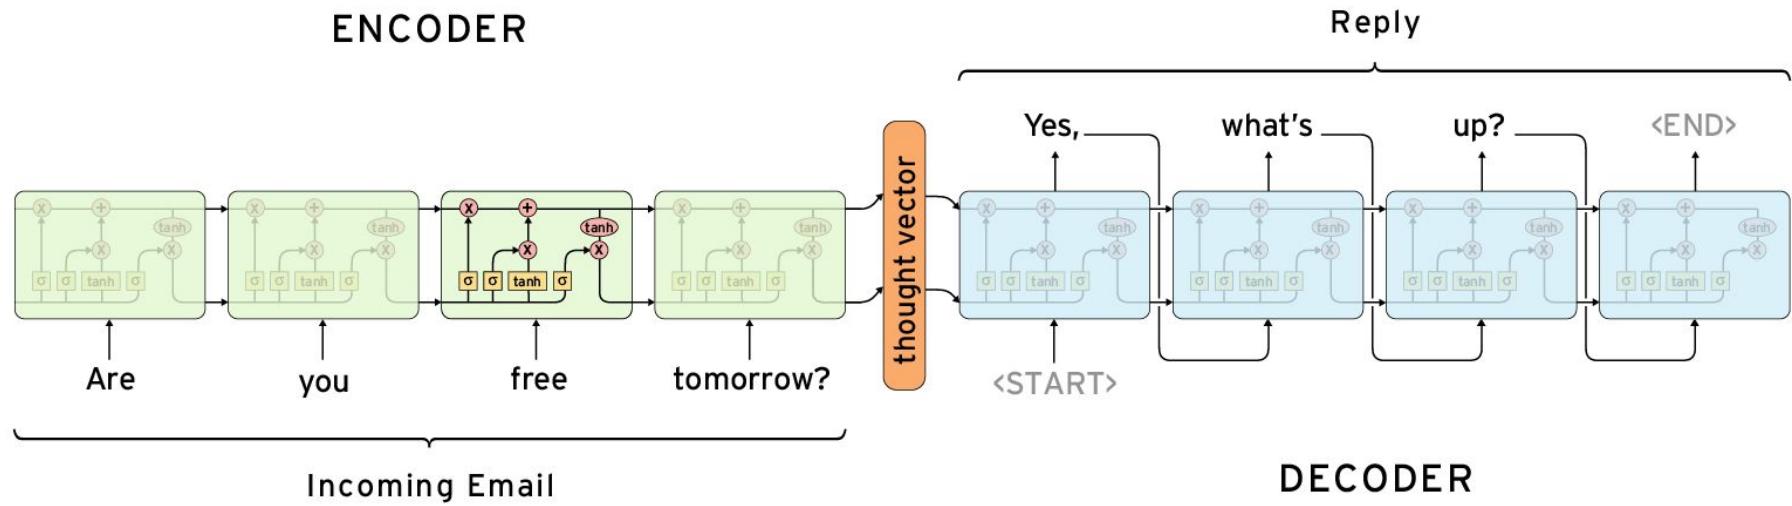
Seq2seq architectures with rnn

Simple encoder, simple decoder

- How to transmit information from encoder to decoder?
 - Initialize hidden state of **decoder** with last hidden state of **encoder**
 - Concatenate last hidden state of **encoder** to each input of **decoder**



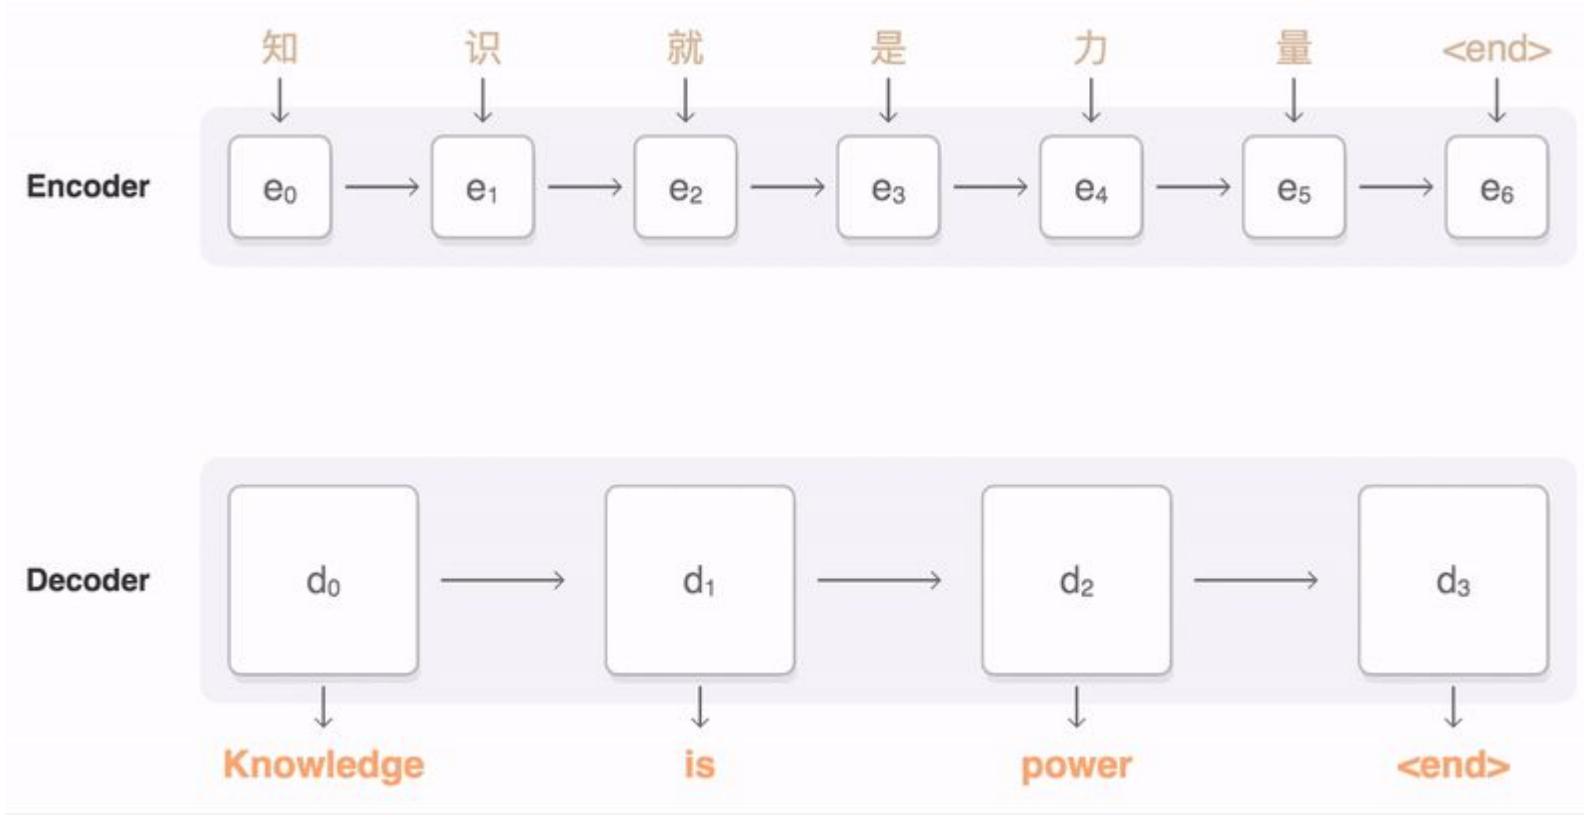
Simple encoder-decoder



Training

1. $\mathbf{X} :=$ last **hidden state of encoder**
2. Init **decoder hidden state** with \mathbf{X}
3. Input <START> token to decoder to start decoding
4. Calculate loss for the decoder output and right output (e.g. cross-entropy)
5. If (**teacher_forcing**):
 - feeding ***right_prev_token*** into decoder
- else:
 - feeding ***decoder_last_output*** into decoder

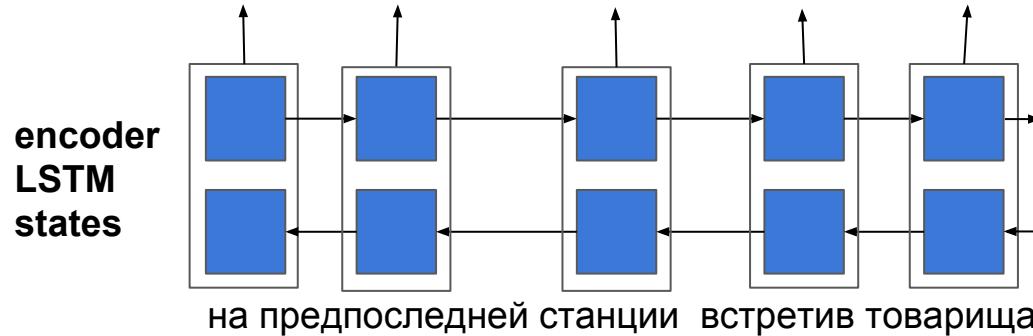
Attention in seq2seq



Outline

1. Seq2seq architectures with rnn
 - a. Applications
 - b. Simple encoder, simple decoder
 - c. **Decoder with attention**
2. Transformer
 - a. Encoder
 - i. Multi-head self-attention
 - ii. LayerNorm & residual connections
 - iii. Position-wise feed-forward
 - iv. Positional Encoding
 - b. Decoder
 - i. Multi-head attention with encoder outputs
 - ii. Masking

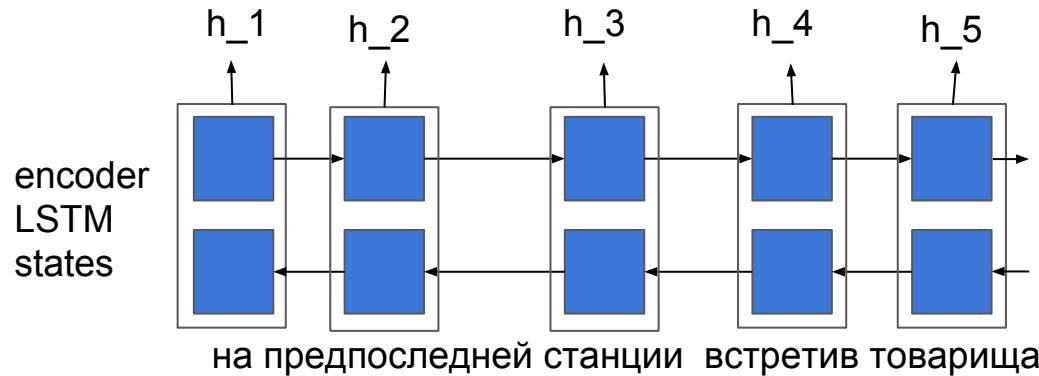
Attention for neural machine translation



Bahdanau et al. "Neural Machine Translation by Jointly Learning to Align and Translate", 2014

Slide credit: Michael Figurnov, DeepBayes school

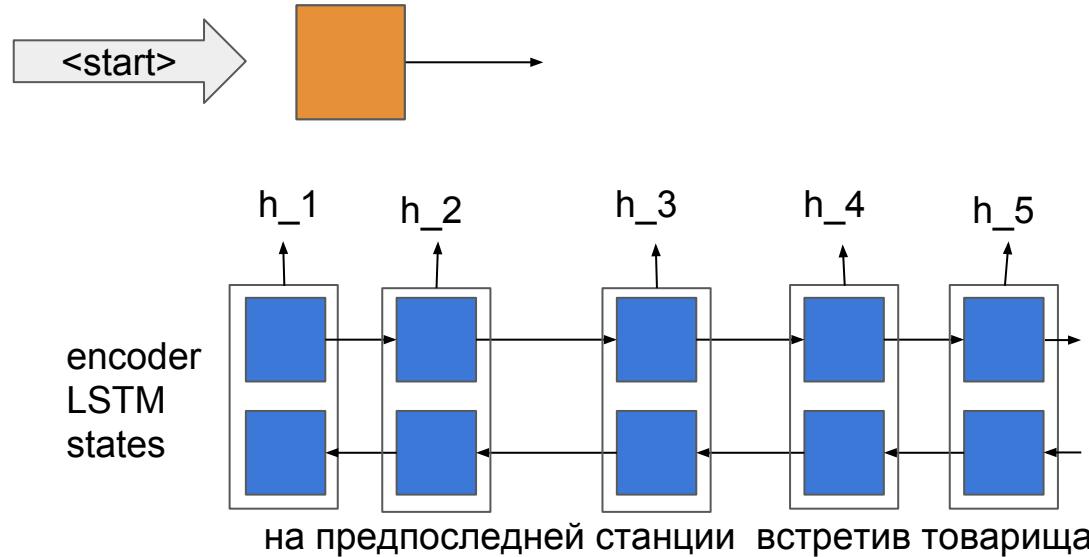
Attention for neural machine translation



Bahdanau et al. "Neural Machine Translation by Jointly Learning to Align and Translate", 2014

Slide credit: Michael Figurnov, DeepBayes school

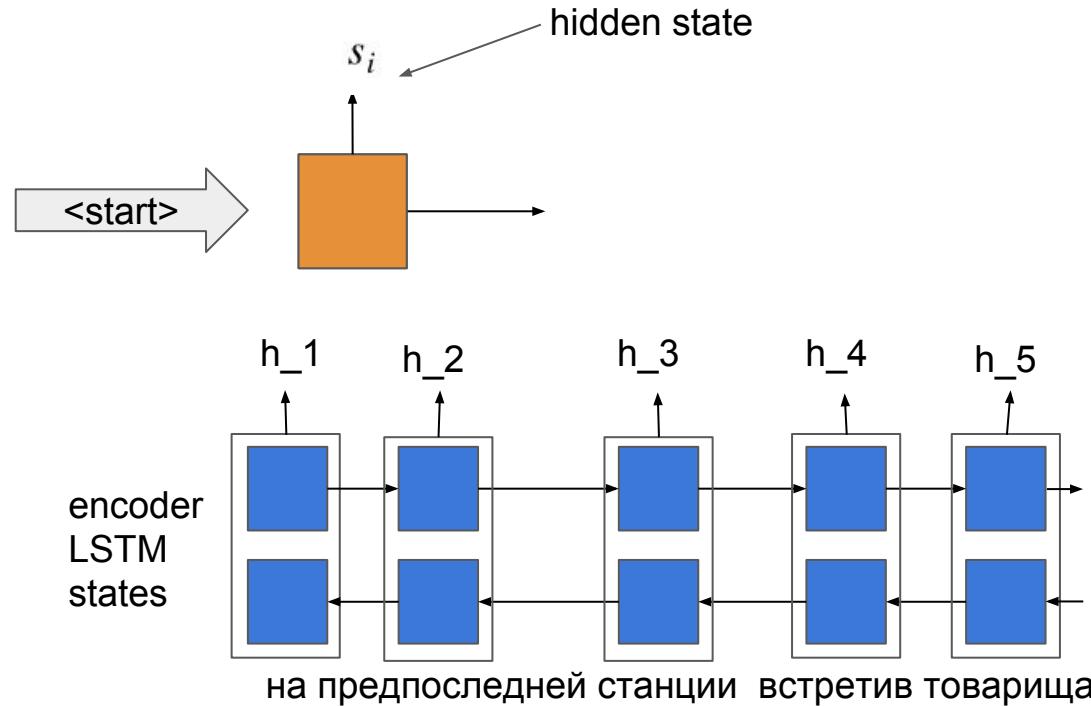
Attention for neural machine translation



Bahdanau et al. "Neural Machine Translation by Jointly Learning to Align and Translate", 2014

Slide credit: Michael Figurnov, DeepBayes school

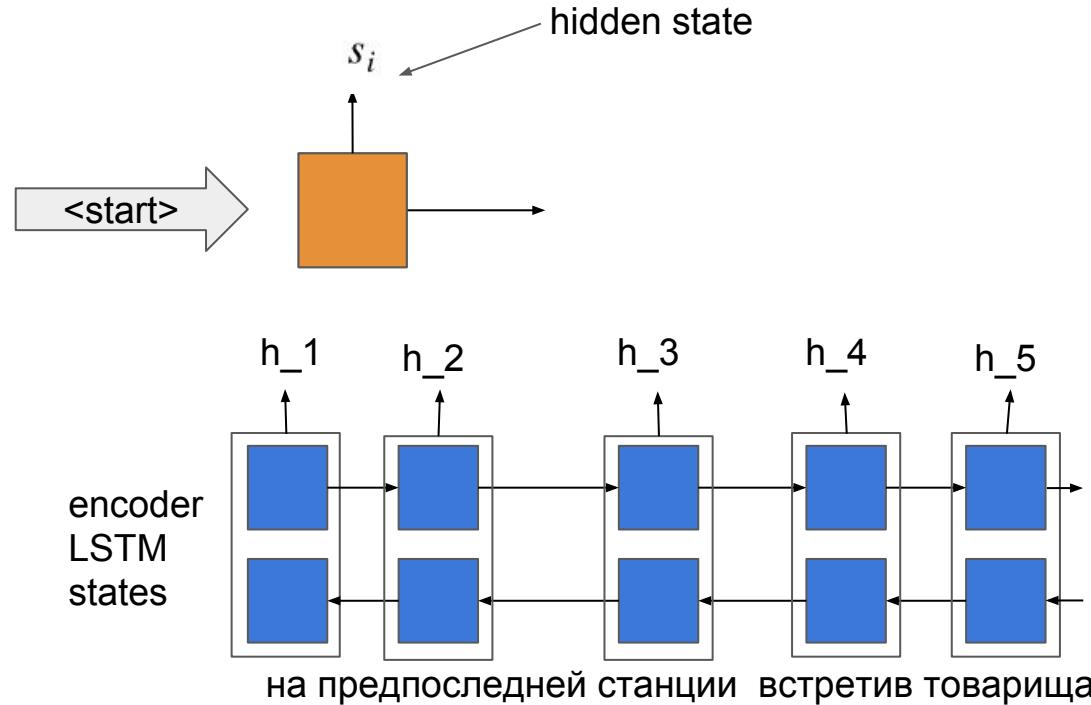
Attention for neural machine translation



Bahdanau et al. "Neural Machine Translation by Jointly Learning to Align and Translate", 2014

Slide credit: Michael Figurnov, DeepBayes school

Attention for neural machine translation



Energy

$$e_{ij} = a(s_{i-1}, h_j)$$

a - some function

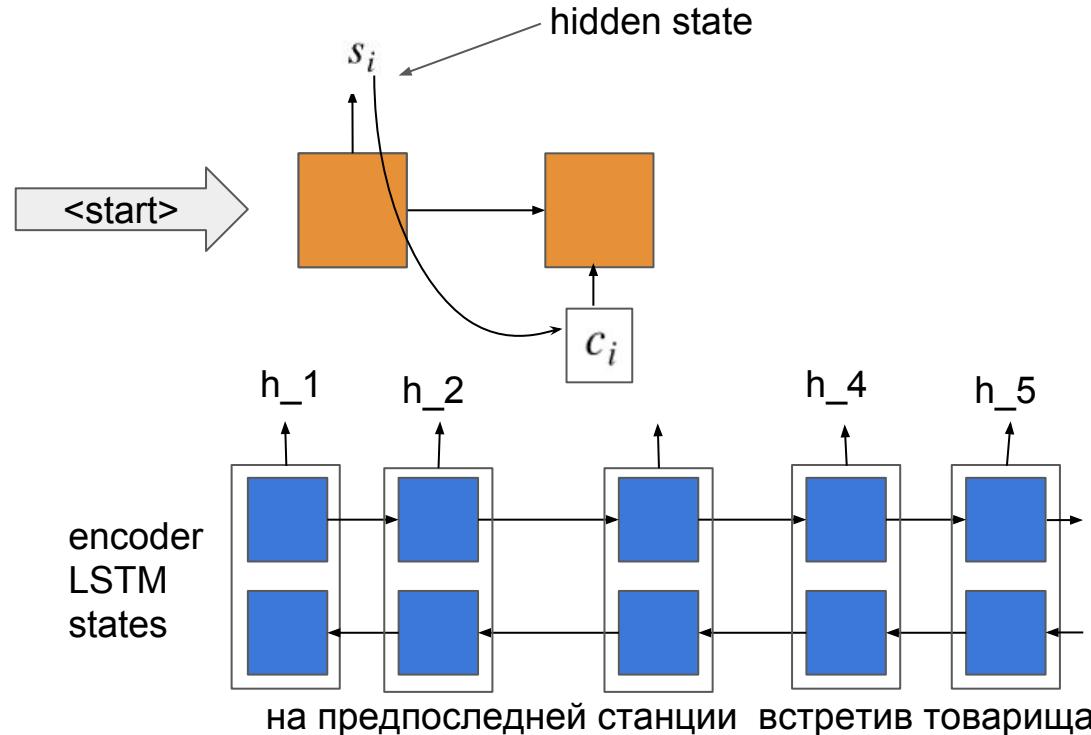
Attention weights

$$a_{ij} = \frac{\exp(e_{ij})}{\sum_{k=1}^T \exp(e_{ik})}$$

Context vector

$$c_i = \sum_{j=1}^{T_x} a_{ij} h_j$$

Attention for neural machine translation



Energy

$$e_{ij} = a(s_{i-1}, h_j)$$

a - some function

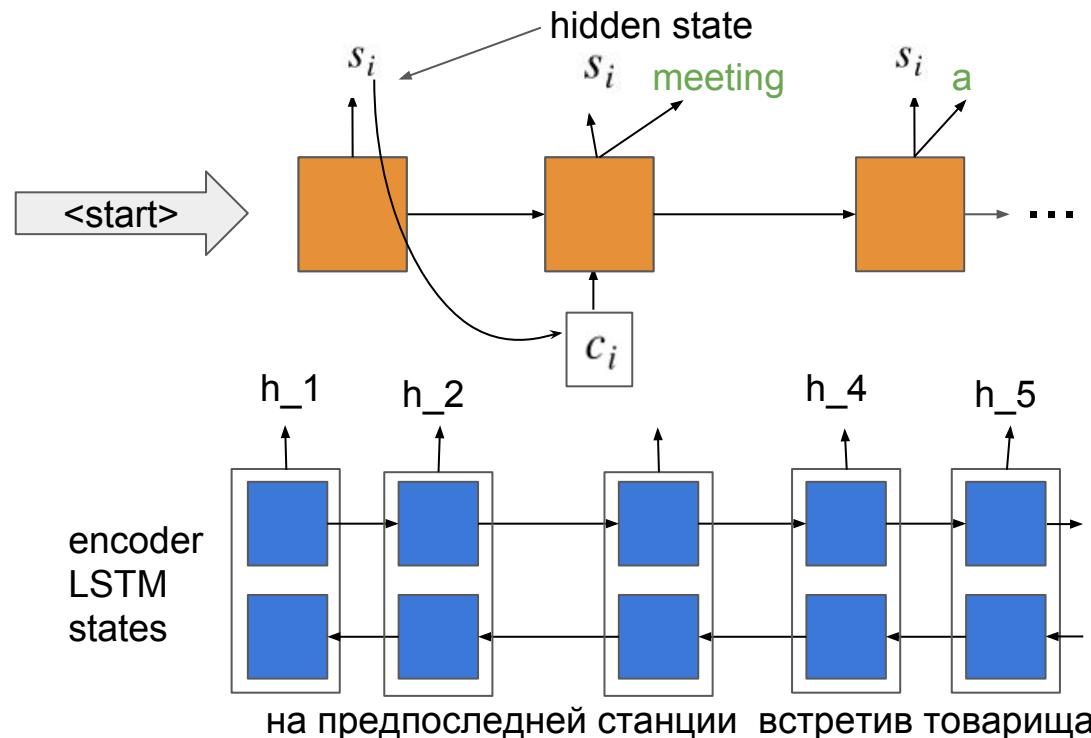
Attention weights

$$a_{ij} = \frac{\exp(e_{ij})}{\sum_{k=1}^T \exp(e_{ik})}$$

Context vector

$$c_i = \sum_{j=1}^{T_x} a_{ij} h_j$$

Attention for neural machine translation



Energy

$$e_{ij} = a(s_{i-1}, h_j)$$

a - some function

Attention weights

$$a_{ij} = \frac{\exp(e_{ij})}{\sum_{k=1}^T \exp(e_{ik})}$$

Context vector

$$c_i = \sum_{j=1}^{T_x} a_{ij} h_j$$

Many ways to compute the energy

$$\text{score}(h, f) = \text{NN}(h, f) = w_3^T \tanh(W_1 h + W_2 f) \quad [1]$$

$$\text{score}(h, f) = h^T f \quad [2]$$

$$\text{score}(h, f) = h^T f / \sqrt{d}, \quad d = \dim(h) \quad [3]$$

additive attention
 $\dim(h) \neq \dim(f)$

} multiplicative
attention
 $\dim(h) = \dim(f)$

[1] Bahdanau et al. "Neural Machine Translation by Jointly Learning to Align and Translate", 2014

[2] Luong et al. "Effective approaches to attention-based neural machine translation", 2015

[3] Vaswani et al. "Attention Is All You Need", 2017

Outline

1. Seq2seq architectures with rnn

- a. Applications
- b. Simple encoder, simple decoder
- c. Decoder with attention

2. Transformer

- a. Encoder
 - i. Multi-head self-attention
 - ii. LayerNorm & residual connections
 - iii. Position-wise feed-forward
 - iv. Positional Encoding
- b. Decoder
 - i. Multi-head attention with encoder outputs
 - ii. Masking

Transformer

NIPS 2017

Attention Is All You Need

Ashish Vaswani*
Google Brain
avaswani@google.com

Noam Shazeer*
Google Brain
noam@google.com

Niki Parmar*
Google Research
nikip@google.com

Jakob Uszkoreit*
Google Research
usz@google.com

Llion Jones*
Google Research
llion@google.com

Aidan N. Gomez* †
University of Toronto
aidan@cs.toronto.edu

Lukasz Kaiser*
Google Brain
lukaszkaiser@google.com

Ilia Polosukhin* ‡
illia.polosukhin@gmail.com

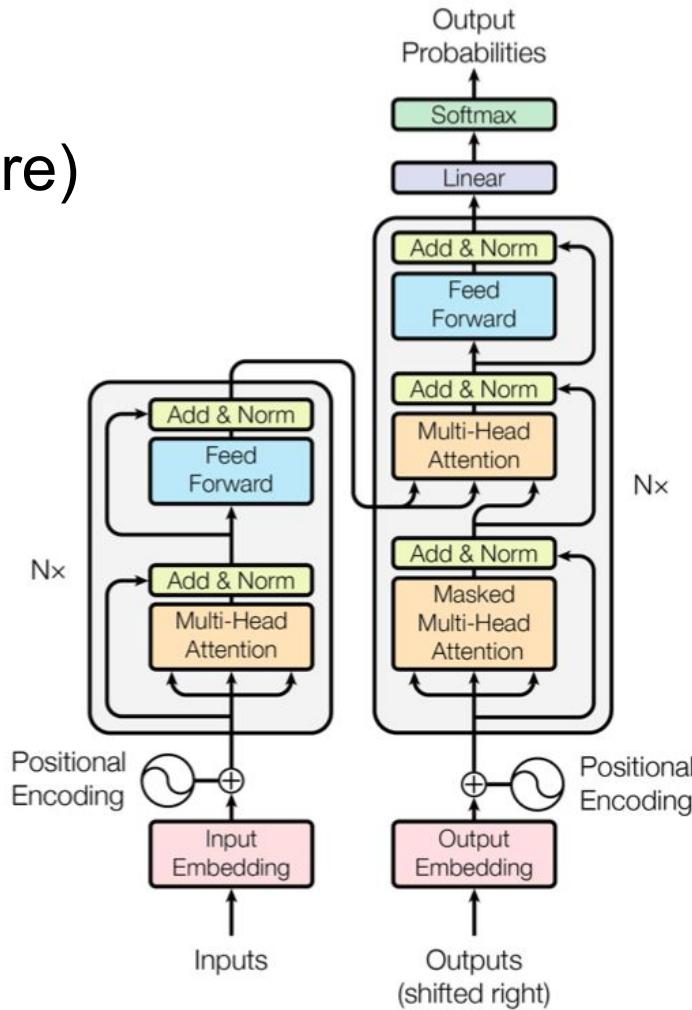
Abstract

The dominant sequence transduction models are based on complex recurrent or convolutional neural networks that include an encoder and a decoder. The best performing models also connect the encoder and decoder through an attention mechanism. We propose a new simple network architecture, the Transformer, based solely on attention mechanisms, dispensing with recurrence and convolutions entirely. Experiments on two machine translation tasks show these models to be superior in quality while being more parallelizable and requiring significantly less time to train. Our model achieves 28.4 BLEU on the WMT 2014 English-to-German translation task, improving over the existing best results, including ensembles, by over 2 BLEU. On the WMT 2014 English-to-French translation task, our model establishes a new single-model state-of-the-art BLEU score of 41.8 after training for 3.5 days on eight GPUs, a small fraction of the training costs of the best models from the literature. We show that the Transformer generalizes well to other tasks by applying it successfully to English constituency parsing both with large and limited training data.

link: <https://papers.nips.cc/paper/7181-attention-is-all-you-need.pdf>

Transformer (model architecture)

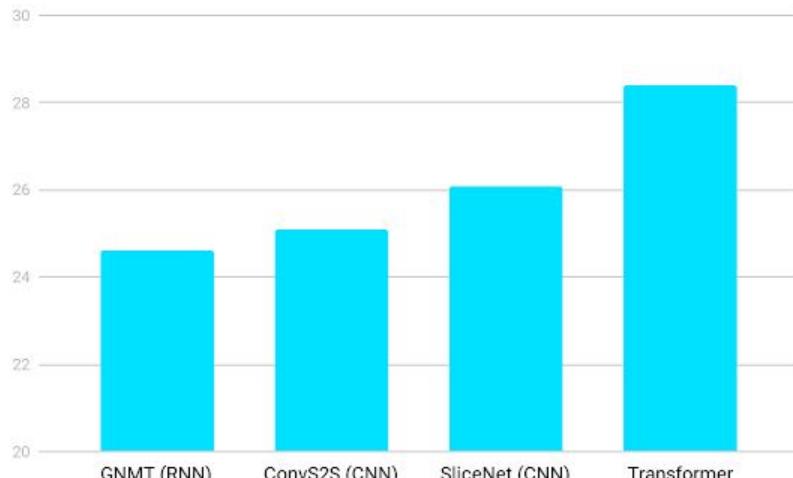
$N = 6$



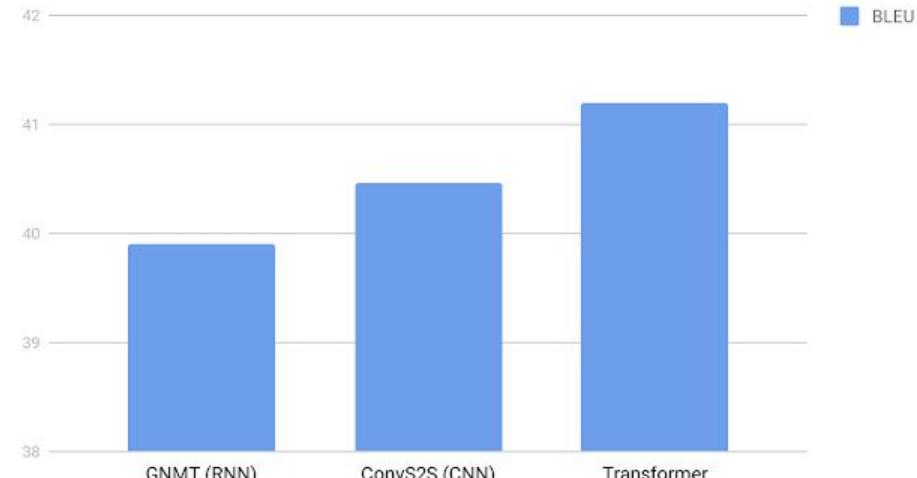
Transformer

- Better BLUE score

English German Translation quality



English French Translation Quality



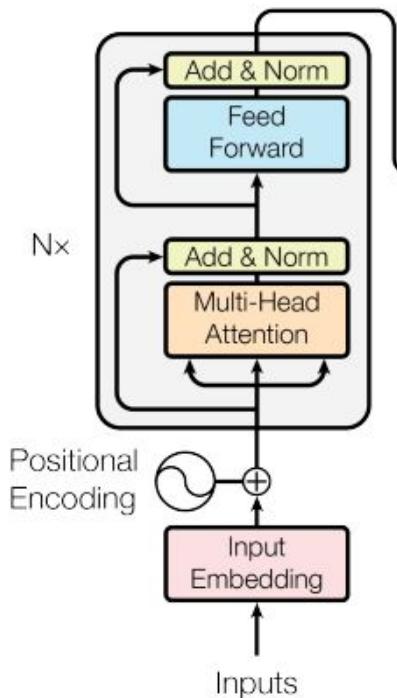
<https://research.googleblog.com/2017/08/transformer-novel-neural-network.html>

GLUE Benchmarks

Rank	Name	Model	URL	Score	CoLA	SST-2	MRPC	STS-B	QQP	MNLI-m	MNLI-mm	QNLI	RTE	WNLI	AX	
1	GLUE Human Baselines	GLUE Human Baselines		87.1	66.4	97.8	86.3/80.8	92.7/92.6	59.5/80.4	92.0	92.8	91.2	93.6	95.9	-	
+	2	Microsoft D365 AI & MSR AI MT-DNN++ (BigBird)		83.8	65.4	95.6	91.1/88.2	89.6/89.0	72.7/89.6	87.9	87.4	95.8	85.1	65.1	41.9	
+	3	王玮	ALICE large (Alibaba DAMO NLP)	83.3	63.5	95.2	91.8/89.0	89.8/88.8	74.0/90.4	87.9	87.4	95.7	80.9	65.1	40.7	
4	Stanford Hazy Research	Snorkel MeTaL		83.2	63.8	96.2	91.5/88.5	90.1/89.7	73.1/89.9	87.6	87.2	93.9	80.9	65.1	39.9	
5	Anonymous	Anonymous BAM		82.3	61.5	95.2	91.3/88.3	88.6/87.9	72.5/89.7	86.6	85.8	93.1	80.4	65.1	40.7	
6	张倬胜	SemBERT		82.0	62.3	94.6	90.6/87.3	87.8/86.7	72.8/89.8	87.2	86.3	93.2	78.5	65.1	40.5	
+	7	Jason Phang	BERT on STILTs		82.0	62.1	94.3	90.2/86.6	88.7/88.3	71.9/89.4	86.4	85.6	92.7	80.1	65.1	28.3
+	8	Jacob Devlin	BERT: 24-layers, 16-heads, 1024-hic		80.5	60.5	94.9	89.3/85.4	87.6/86.5	72.1/89.3	86.7	85.9	92.7	70.1	65.1	39.6
9	Neil Houlsby	BERT + Single-task Adapters		80.2	59.2	94.3	88.7/84.3	87.3/86.1	71.5/89.4	85.4	85.0	92.4	71.6	65.1	9.2	
10	Alec Radford	Singletask Pretrain Transformer		72.8	45.4	91.3	82.3/75.7	82.0/80.0	70.3/88.5	82.1	81.4	87.4	56.0	53.4	29.8	
11	GLUE Baselines	BiLSTM+ELMo+Attn		70.0	33.6	90.4	84.4/78.0	74.2/72.3	63.1/84.3	74.1	74.5	79.8	58.9	65.1	21.7	
		BiLSTM+ELMo		67.7	32.1	89.3	84.7/78.0	70.3/67.8	61.1/82.6	67.2	67.9	75.5	57.4	65.1	21.3	
		Single Task BiLSTM+ELMo+Attn		66.5	35.0	90.2	80.2/68.8	55.5/52.5	66.1/86.5	76.9	76.7	76.7	50.3	65.1	27.9	
		Single Task BiLSTM+ELMo		66.4	35.0	90.2	80.8/69.0	64.0/60.2	65.6/85.7	72.9	73.4	71.7	50.1	65.1	19.5	
		GenSen		66.1	7.7	83.1	83.0/76.6	79.3/79.2	59.8/82.9	71.4	71.3	78.6	59.2	65.1	20.6	
		BiLSTM+Attn		65.6	18.6	83.0	83.9/76.2	72.8/70.5	60.1/82.4	67.6	68.3	74.3	58.4	65.1	17.8	
		BiLSTM		64.2	11.6	82.8	81.8/74.3	70.3/67.8	62.5/84.2	65.6	66.1	74.6	57.4	65.1	20.3	

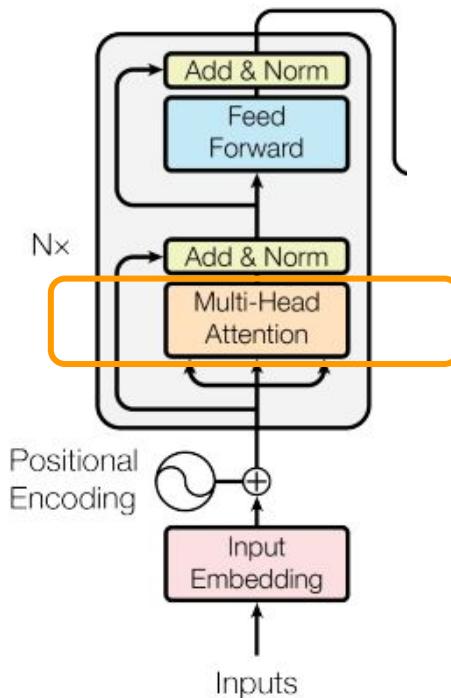
<https://gluebenchmark.com/leaderboard>

Encoder



- Multi-head self-attention
- LayerNorm & Residual connections
- Position-wise feed-forward
- Positional Encoding

Encoder

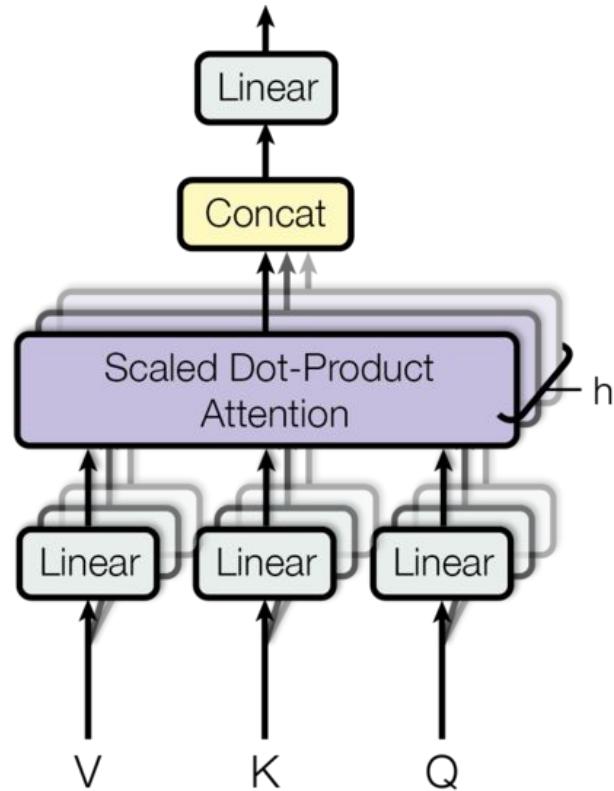


- Multi-head self-attention
- LayerNorm & Residual connections
- Position-wise feed-forward
- Positional Encoding

Multi-head self-attention

How does one head work?

1. Value (V), Key (K), Query (Q) -- наборы векторов слов.
(Value & Key одни и те же входные вектора, но подаются они на разные входы независимо)
2. Преобразуем входные V, K, Q каждый своим отдельным обучаемым линейным преобразованием.
3. Считаем скалярное произведение каждого вектора из Q с каждым вектором из K (*т.е. величину схожести каждого слова с каждым*). Полученные произведения делим на корень из размерности вектора.
4. Для конкретного вектора q из Q мы получим преобразованный вектор, путем сложения всех векторов из V, с нормализованными весами, полученными в п. 3 с этим самым вектором q .
5. Возвращаем набор преобразованных векторов q.



Multi-head self-attention

Intuition

Keys = Values = Queries =
embeddings(["In", "my", "humble", "opinion"])

$q := \text{In};$

$q_{\text{out}} := \text{my} * (\text{In}, \text{my}) + \text{humble} * (\text{In}, \text{humble}) + \text{opinion} * (\text{In}, \text{opinion})$

Multi-head self-attention

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O$$

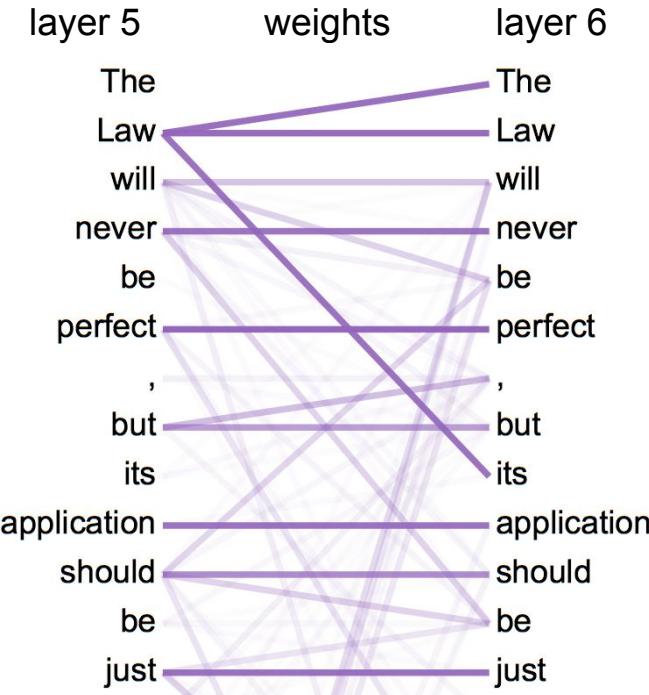
$$\text{where } \text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$$

Multi-head attention allows the model to jointly attend to information from different representation subspaces at different positions.

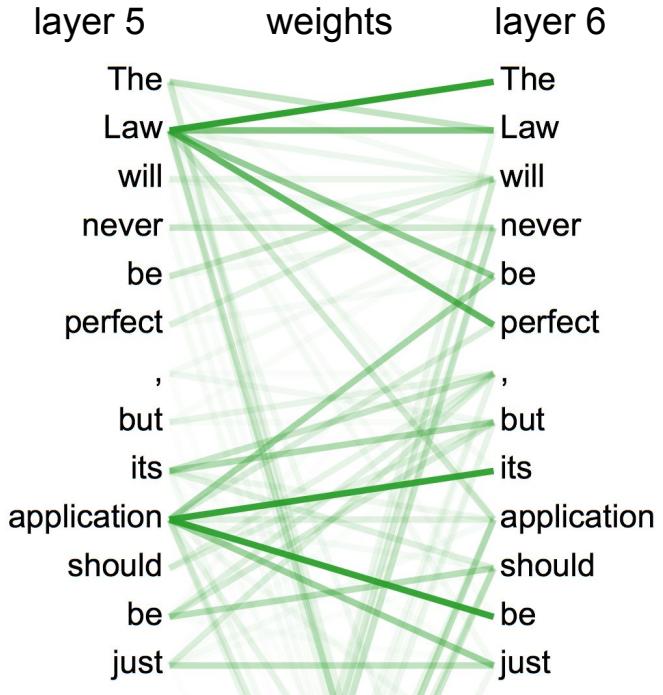
Dimension of the returned single matrix after applying Multi-head sublayer is the same as the dimension of any its input.

What do the self-attention heads look at?

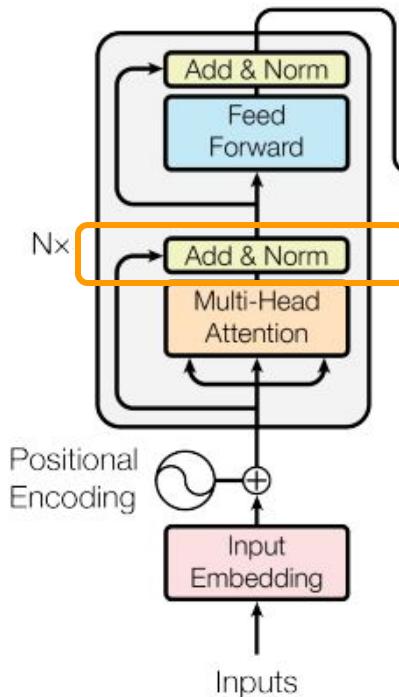
attention head #1



attention head #2

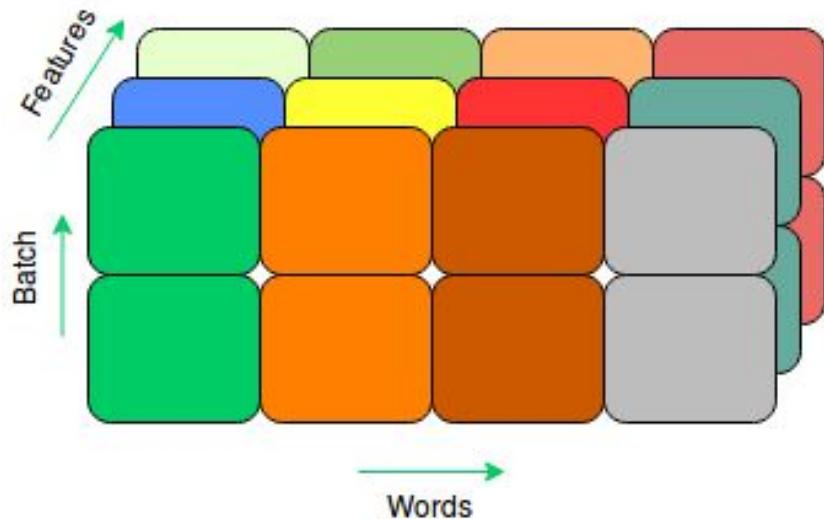


Encoder

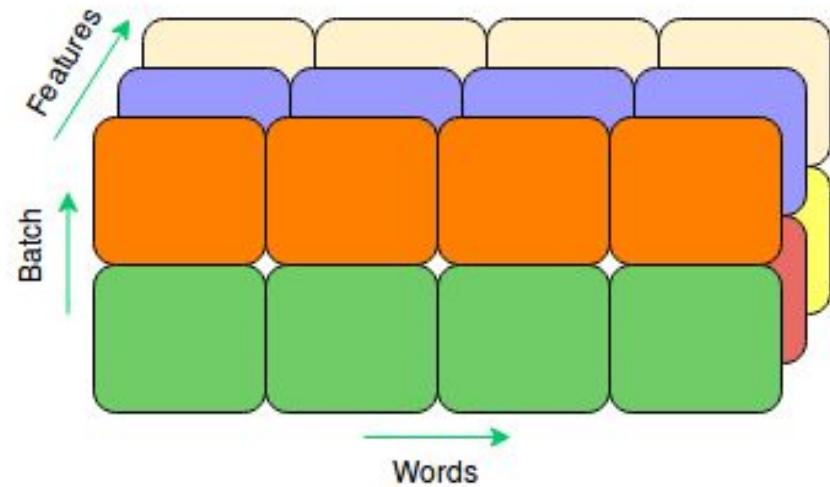


- Multi-head self-attention
- LayerNorm & Residual connections
- Position-wise feed-forward
- Positional Encoding

LayerNorm & Residual connections



Batch Norm



Layer Norm

LayerNorm & Residual connections

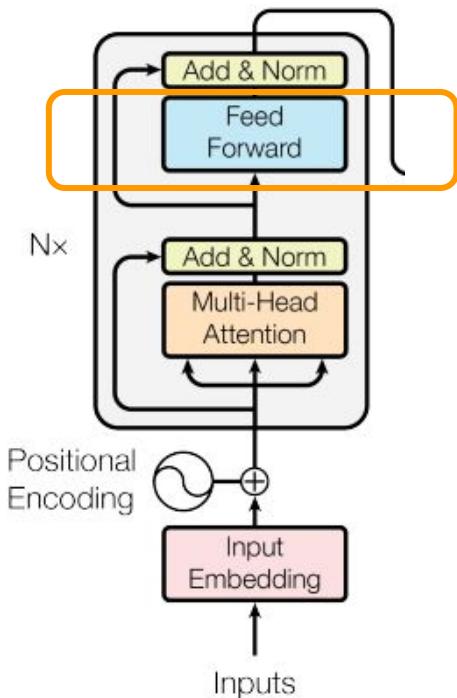
```
class LayerNorm(nn.Module):
    "Construct a layernorm module (See citation for details)."
    def __init__(self, features, eps=1e-6):
        super(LayerNorm, self).__init__()
        self.a_2 = nn.Parameter(torch.ones(features))
        self.b_2 = nn.Parameter(torch.zeros(features))
        self.eps = eps

    def forward(self, x):
        mean = x.mean(-1, keepdim=True)
        std = x.std(-1, keepdim=True)
        return self.a_2 * (x - mean) / (std + self.eps) + self.b_2
```

LayerNorm($x + \text{Sublayer}(x)$)

Residual connection

Encoder



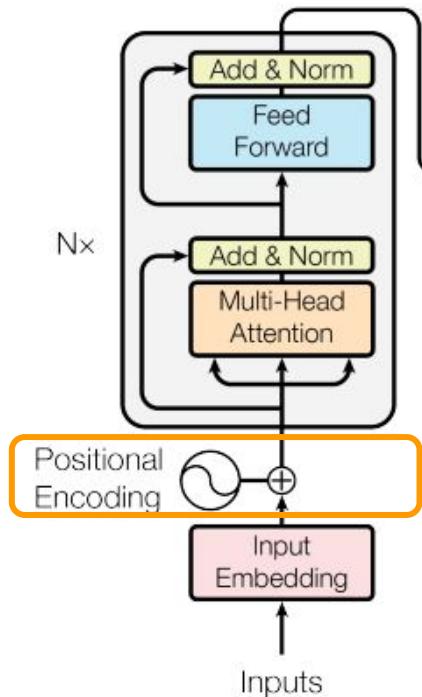
- Multi-head self-attention
- LayerNorm & Residual connections
- Position-wise feed-forward
- Positional Encoding

Position-wise feed-forward

$$\text{FFN}(x) = \max(0, xW_1 + b_1)W_2 + b_2$$

- ReLu + two dense layers
- $\text{dim}(\text{input}) = \text{dim}(\text{output})$

Encoder



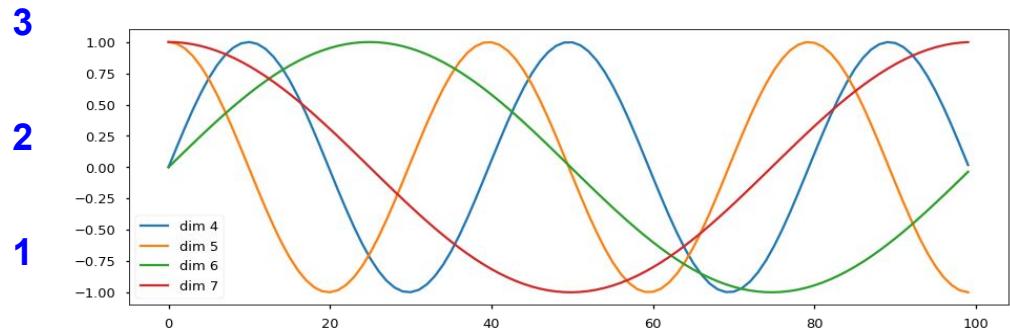
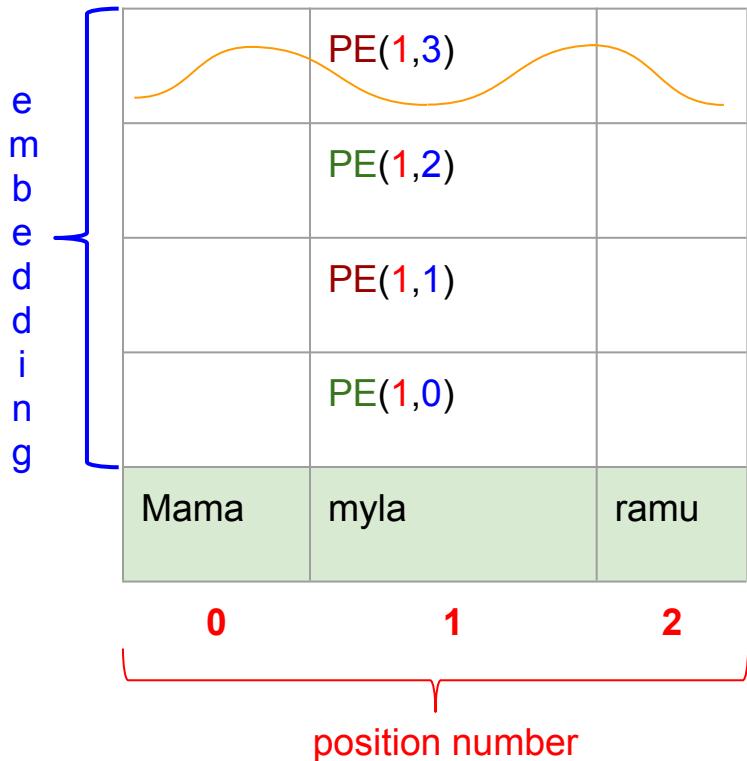
- Multi-head self-attention
- LayerNorm & Residual connections
- Position-wise feed-forward
- **Positional Encoding**

Positional encoding

- Learning embeddings for each position
 - (-) Bad embeddings for high-numbered positions
 - (-) More parameters needed
 - (-) Length restrictions
- Using sinusoids to encode position
 - **adding position vector to embedding vector**
 - concatenating position vector to embedding vector

Same scores
with these two
approaches

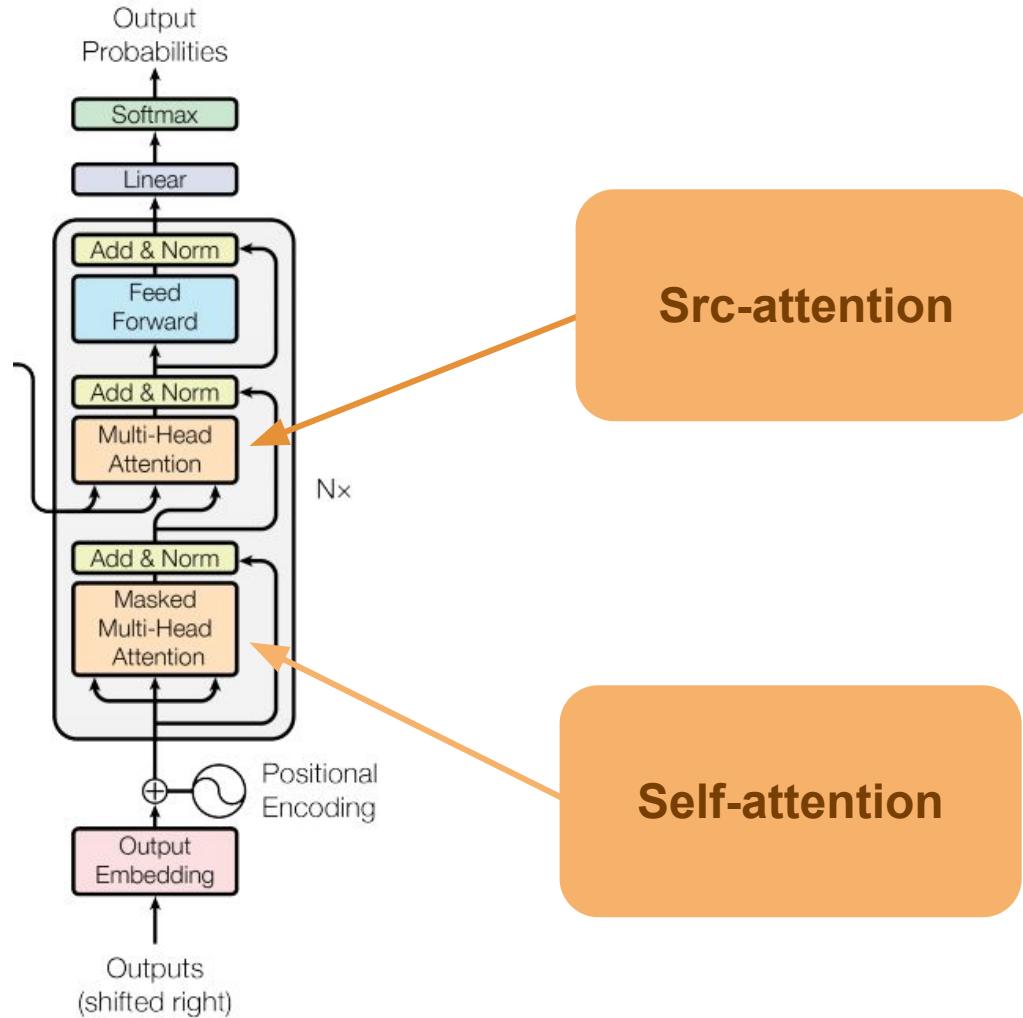
Positional encoding



$$PE_{(pos,2i)} = \sin(pos/10000^{2i/d_{\text{model}}})$$

$$PE_{(pos,2i+1)} = \cos(pos/10000^{2i/d_{\text{model}}})$$

Decoder



Decoder

```
class DecoderLayer(nn.Module):
    "Decoder is made of self-attn, src-attn, and feed forward (defined below)"
    def __init__(self, size, self_attn, src_attn, feed_forward, dropout):
        super(DecoderLayer, self).__init__()
        self.size = size
        self.self_attn = self_attn
        self.src_attn = src_attn
        self.feed_forward = feed_forward
        self.sublayer = clones(SublayerConnection(size, dropout), 3)

    def forward(self, x, memory, src_mask, tgt_mask):
        "Follow Figure 1 (right) for connections."
        m = memory
        x = self.sublayer[0](x, lambda x: self.self_attn(x, x, x, tgt_mask))
        x = self.sublayer[1](x, lambda x: self.src_attn(x, m, m, src_mask))
        return self.sublayer[2](x, self.feed_forward)
```

<https://research.googleblog.com/2017/08/transformer-novel-neural-network.html>

Optional part

- Regularization
(Label smoothing)
- BPE —
Byte Pair Encoding for NMT
- Beam Search



Regularization

Label smoothing (LRS)

$$q'(k|x) = (1 - \epsilon)\delta_{k,y} + \epsilon u(k)$$

- Replace true label distribution $q(k|x) = \delta_{k,y}$ with $q'(k|x)$
- $u(k)$ - independent from x distribution (e.g. uniform)
- ϵ - just a small real number
- $\delta_{k,y}$ - Dirac delta; equals 1 if $k = y$ else equals 0;

This hurts perplexity, as the model learns to be more unsure, but improves accuracy and BLEU score.

BPE

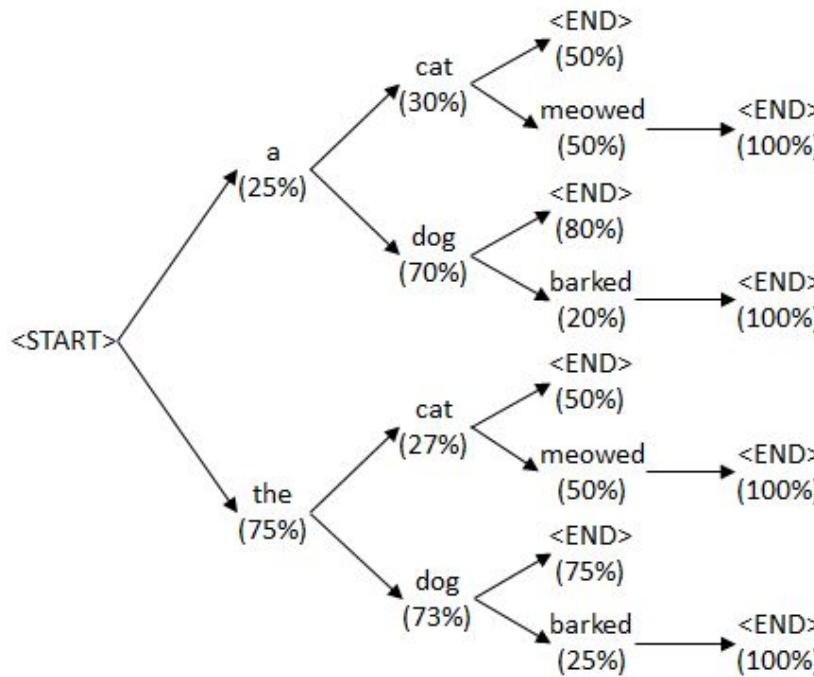
* **Byte Pair Encoding (BPE)** (Gage, 1994) is a simple data compression technique that iteratively replaces the most frequent pair of bytes in a sequence with a single, unused byte. We adapt this algorithm for word segmentation. Instead of merging frequent pairs of bytes, we merge characters or character sequences.

- Official implementation:
<https://github.com/rsennrich/subword-nmt>

Example

- Input:
 - “It is the case of Alexander Nikitin.”
- After applying BPE:
 - It is the case of Alexander Ni@® ki@® tin .

Beam Search



PyTorch OpenNMT implementation:

<https://github.com/OpenNMT/OpenNMT-py/blob/master/onmt/translate/Beam.py>

The background of the image is a deep space scene featuring a large, luminous nebula with swirling patterns of orange, yellow, blue, and white. A single, very bright blue star is visible in the upper right corner. The overall texture is grainy and has a cinematic feel.

вопросы